

# Пишем load balancer на Spring

**Sergey Morenets** 

October,  $13^{th} - 14^{th}$  2018
• Sergey Morenets, 2018



#### **DEVELOPER 12 YEARS**



TRAINER 4 YEARS

WRITER 3 BOOKS



#### **FOUNDER**







#### **SPEAKER**























#### Goals



Completed project

**Practice** 

Load balancing principles

• Serge', 11101011010, 2010

#### **Agenda**

DISCOVERY

- ✓ REST and REST services
- ✓ Spring and Spring Boot usage
- ✓ Client-side and server-side load balancers
- ✓ Algorithms and principles
- ✓ Functional testing
- ✓ Monitoring using Spring Boot Actuator
- Error handling
- Configuration and customization
- Performance and security testing
- ✓ Reactive programming







Talk is cheap. Show me the code.

— Linus Torvalds —

AZ QUOTES

#### Task 1. Reviewing project



- Import load-balancer-rest project into your IDE (you should import it as Maven or Gradle project) and open book-service sub-project.
- 2. Review project structure, dependencies and REST services.
- 3. Run **BookApplication** and try to invoke REST services using Postman application.
- 4. Review application logs and Spring configuration



## **Load balancing**





#### **Load balancing**



- Distributing incoming network traffic through group of backend servers (server pool)
- Client request router
- ✓ Allows horizontal scaling, optimize resource use, maximizes speed and capacity utilization
- Automatic configuration and error handling
- Health checking and failovers
- DDos attack protection
- Client authentication and firewall
- Caching

#### **Load balancers**











#### **Load balancers**







#### Load balancer workflow

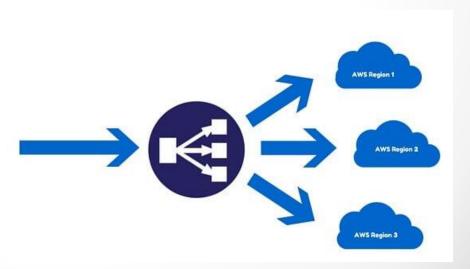


- ✓ Configure back-end servers (static or dynamic)
- Health checks and monitoring
- ✓ Choose back-end server to route request
- Session persistence
- Error handling
- Elastic load balancing

#### **Load balancing algorithms**



- ✓ Round Robin
- ✓ Least utilization
- Geographic location



#### Task 2. Order service



- 1. Review order-service sub-project and its dependencies.
- Review OrderController class and complete its makeOrder REST service. This service should query book-service and gets book by id.
- 3. Run both services and make sure they communicate properly. What will happen if book-service shuts down?
- 4. Write integration (or unit-test) for OrderController.



#### **Mapping annotations**



```
@Bean
public RestTemplate restTemplate(RestTemplateBuilder builder) {
    return builder.build();
}
```

```
@Autowired
private RestTemplate restTemplate;
```

Spring bean

Spring configuration

### RestTemplateBuilder



Method	Description
rootUri()	Applies root URI for each request
Interceptors()	Applies client request interceptors
errorHandler()	Specifies response error handler
basicAuthorization()	Specifiies username/password for HTTP basis authorization
customizers()	Sets customizers to change output RestTemplate
setConnectionTimeout()	Change connection timeout
setReadTimeout()	Changes read timeout

#### **Testing REST services. Configuration**



```
@SpringJUnitWebConfig(DemoApplication.class)
                                        Loads Spring context
@SpringBootTest
@AutoConfigureWebClient Configure RestTemplateBuilder
                                     Configure JacksonTester
@JsonTest
public class OrderControllerTest {
    @Autowired
    private RestTemplate restTemplate;
    @Autowired
    private OrderController orderController;
    @Autowired
    private JacksonTester<Book> bookTester;
    private MockRestServiceServer mockServer;
```

#### **Testing REST services. Tests**



```
Configure MockServer
@BeforeEach
public void setUp()
    mockServer = MockRestServiceServer.createServer(restTemplate);
@Test
public void makeOrder validBookId orderCreated() throws IOException
    Book book = new Book();
    book.setId(1);
    mockServer.expect(requestTo("http://localhost:8081/books"))
        .andRespond(withSuccess(bookTester.write(book).getJson(),
                MediaType.APPLICATION_JSON_UTF8));
    Order order = orderController.makeOrder(1);
                                                    Mock response
    mockServer.verify();
    assertEquals(order.getBookId(), 1);
```

#### Task 3. Order service performance



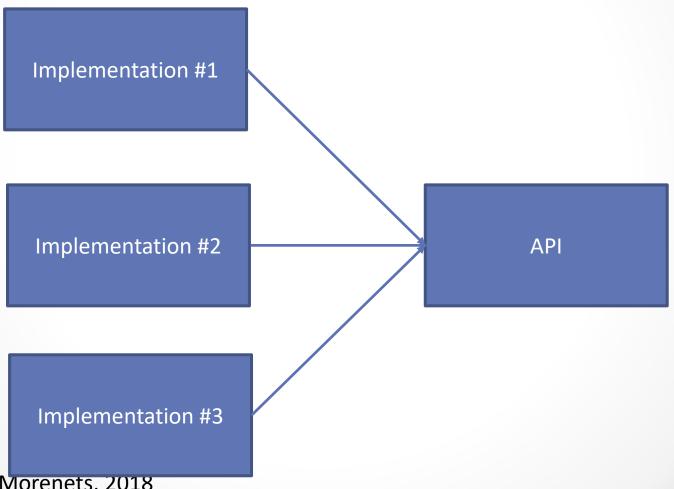
- 1. Download and install Apache JMeter 5.0 (and later), for example, into c:/JMeter folder.
- 2. Run JMeter executable file in c:/Jmeter/bin folder.
- 3. Click on **Test Plan** node and add new Thread Group. Specify number of threads (users), for example, 10, then ramp-up time and loop count (10 or 20 is recommended).
- Click on Thread Group node and add new Sampler -> HTTP request



# Client load balancer vs Server load balancer

#### Client-side load balancer





#### Task 4. Starting balancer component

DISCOVERY

- 1. Review client-balancer sub-project.
- 2. Add Spring Boot startup configuration
- Implement basic interfaces and classes that will be part of load-balancer API
- 4. Add static configuration that will store list of back-end servers.



# Task 5. Link order-service and balancer DISCOVERY

- 1. Include client-balancer dependency to the order-service project.
- Add configuration property in order-service that will store back-end server-list
- Create implementation of the load balancer API that randomly chooses target server and declare Spring bean of this implementation.



#### **Bean Validation API**



Annotation	Description
@NotEmpty	Element should not be empty
@NotBlank	Element should not be blank
@Min	Specifies minimum value of the element
@Max	Specifies maximum value of the element
@NotNull	Element should not be null
@Pattern	Specifies regular expression pattern
@Email	Element should be email
@Future	Element should be time or instant in the future
@Negative	Element should be negative number
@Size	Element size (length) should be between min and max

#### **Enable auto-configuration**



```
@Component
public class CustomService {
public class CustomService {
}
}
```

```
@Configuration
public class CustomServiceAutoConfiguration {
    @Bean
    @ConditionalOnMissingBean
    public CustomService customService() {
        return new CustomService();
    }
}
```

# Task 6. Validation and auto-configuration

- Add validation annotations to verify that configuration data are valid
- 2. Enable auto-configuration for the default implementation of the client-side load balancer



#### **Actuator**



- Helps manage and monitor applications when pushed to production
- ✓ Accessible via HTTP, JMX or remote shell
- ✓ You can't manage what you can't measure



# **Spring Boot Actuator**



- ✓ Series of endpoints to help manage your Spring application
- Reads properties and spring beans and then returns a JSON view
- ✓ Allows direct access to non functional application information without having to open an IDE or a command prompt



#### **Spring Boot Actuator. Maven**



```
<dependency>
     <groupId>org.springframework.boot</groupId>
     <artifactId>spring-boot-actuator</artifactId>
     <version>${spring.boot.version}</version>
</dependency>
```

#### **Spring Boot Actuator. Endpoints**



```
▼ links:
  ▼ self:
                     "http://localhost:8080/actuator"
      href:
      templated:
                     false
  ▼ health:
      href:
                     "http://localhost:8080/actuator/health"
                     false
      templated:
  ▼ info:
      href:
                     "http://localhost:8080/actuator/info"
      templated:
                     false
```

https://localhost:8080/actuator

# Spring Boot Actuator. Endpoints

Endpoint	Description
/actuator/beans	Displays list of Spring beans in the application
/actuator/metrics	Shows application metrics
/actuator/env	Exposes environment variables
/actuator/loggers	Allows to read/change logger settings
/actuator/health	Shows application health information
/actuator/mappings	Display a list of @RequestMapping info
/actuator/conditions	Displays auto-configuration report with filtering support
/actuator/info	Displays application-related info
/actuator/scheduledtasks	Display scheduled tasks

Sergey Morenets, 2018

# **Endpoints management**



management.endpoints.web.exposure.include=\*



management.endpoint.loggers.cache.time-to-live=50s

application.properties

## **Health information**



```
{
    "status": "UP" /actuator/health
}
```

Indicators	
CassandraHealthIndicator	Checks that Cassandra server is up
DiskSpaceHealthIndicator	Checks for low disk space.
KafkaHealthIndicator	Checks that Kafka is up.
Neo4jHealthIndicator	Checks that Neo4j server is up
MongoHealthIndicator	Checks that a Mongo database is up.
RabbitHealthIndicator	Checks that a Rabbit server is up
RedisHealthIndicator	Checks that a Redis server is up
SolrHealthIndicator	Checks that a Solr server is up
MailHealthIndicator	Checks that a mail server is up

# **Metrics. Spring Boot 1.5**



/application/metrics

```
"mem":185856,
"mem.free":98617,
"processors":4,
"uptime":36557,
"heap.committed":185856,
"heap.init":131072,
"heap.used":87238,
"heap": 1860608,
"threads.peak":20,
"threads.daemon":17,
"threads":19,
"classes":9122.
"classes.loaded":9122,
"classes.unloaded":0,
"gc.ps_scavenge.count":29,
"gc.ps_scavenge.time":394,
"gc.ps_marksweep.count":4,
"gc.ps_marksweep.time":920
```

## **Metrics. Spring Boot 2.0**



```
names:
                                "jvm.buffer.memory.used"
                        0:
                                "jvm.memory.used"
                        1:
                                "jvm.gc.memory.allocated"
                        2:
                                "jvm.memory.committed"
                        3:
/actuator/metrics
                                "tomcat.sessions.created"
                        4:
                                "tomcat.sessions.expired"
                        5:
                                "tomcat.global.request.max"
                        6:
                                "tomcat.global.error"
                        7:
                                "jvm.gc.max.data.size"
                        8:
                                "logback.events"
                        9:
                                "system.cpu.count"
                        10:
                                "jvm.memory.max"
                        11:
```

# Metrics. Spring Boot 2.0



```
"jvm.memory.used"
name:
measurements:
▼ 0:
     statistic:
                    "VALUE"
     value:
                    106108720
availableTags:
 ▼ 0:
                    "area"
     tag:
   ▼ values:
                    "heap"
        0:
                    "nonheap"
        1:
```

/actuator/metrics/jvm.memory.used

# Metrics management



- Based on Micrometer façade
- ✓ Supports different monitoring solutions:













# Graphite



- ✓ Open-source monitoring system
- ✓ Released in 2008
- Contains UI application (based on Python and Django), Carbon service and Whisper file system
- ✓ Integrated with Diamond, Ganglia, Grafana, Graphen and others



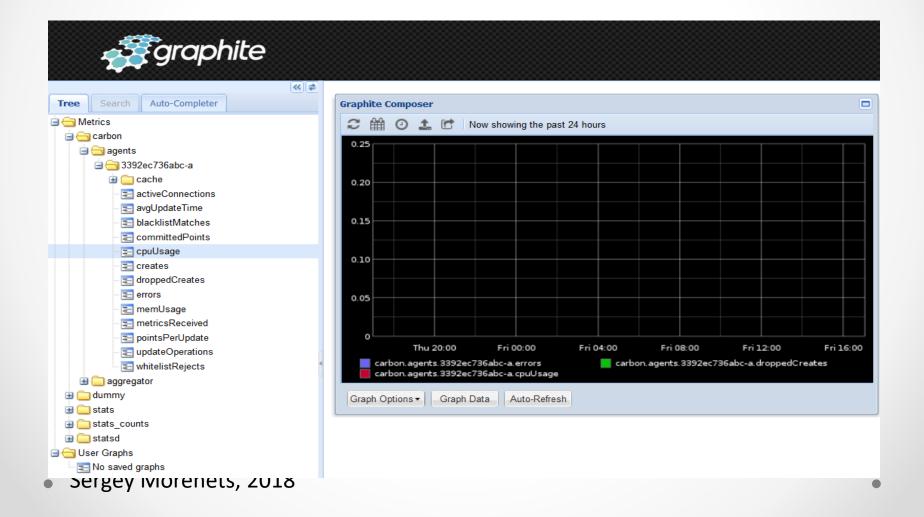
## Micrometer and Graphite



Use default settings to report metrics to Graphite server

# **Graphite Dashboard**





#### **Custom metrics**



```
@Component
                                             Micrometer
public class SampleComponent {
    private final Counter invoiceCounter;
    public SampleComponent(MeterRegistry registry) {
        this.invoiceCounter = registry.counter(
                "total.invoices");
    public void handleInvoice(Invoice invoice) {
        this.invoiceCounter.increment();
```

## Info contributors



Contributors	Description
EnvironmentInfoContributor	Returns all the environment properties for keys started with <b>info</b>
GitInfoContributor	Returns information from a git.properties file
BuildInfoContributor	Returns build information from META-INF/build-info.properties file

# **Application information**



```
info.application.name=Spring Boot
info.application.description=Test application
info.application.version=0.1.0
```

application.properties



```
{
    "application": {
        "version": "0.1.0",
        "description": "Test application",
        "name": "Spring Boot"
    }
}
/actuator/info
```

#### Task 7. Load balancer metrics



- 1. Add new metrics to order-service project that represent statistics of load balancer activities.
- 2. Add Micrometer provider dependency to order-service project, for example, for Graphite
- 3. Start Graphite using Docker image or standalone installation:
- 4. Call REST services of the order-service project and make sure new metrics are displayed in the Graphite dashboard: http://localhost/dashboard

#### **Failsafe library**



- ✓ Lightweight zero-dependency library for handling failures
- ✓ Supports retries, circuit breakers, fallbacks, event listeners asynchronous execution, execution tracking

#### Failsafe library. Examples



## **RetryPolicy class API**



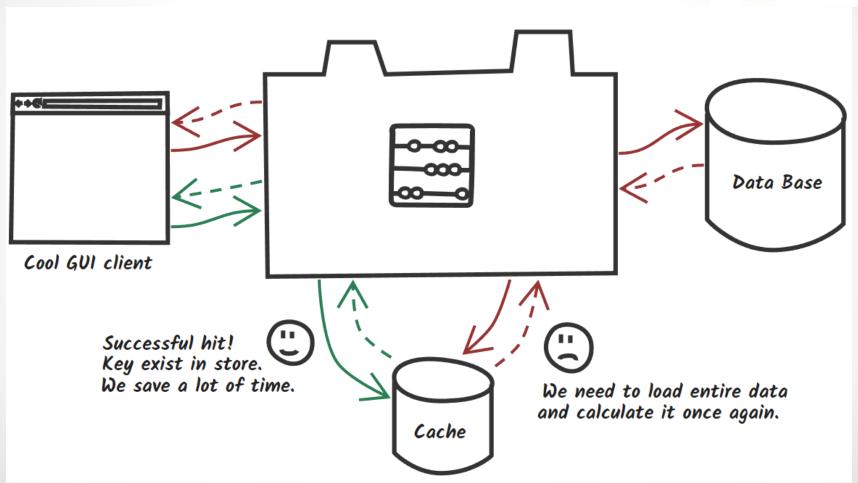
Annotation	Description
abortIf()	Abort execution if specified predicate matches the result
abortOn()	Abort execution if specified exception occurs
retryIf()	Retries if predicate matches the result
retryOn()	Retries attempt if specified exception occurs
withBackoff()	Sets backoff for the next retry
withDelay()	Sets delay between retries
withMaxDuration()	Specifies max duration for the overall retries
withMaxRetries()	Specifies maximum number of retries

# Task 8. Health check and retry policy DISCOVERY

- Create RetryPolicy interface that will define how load balancer act if request to the server fails. What should be default retry policy?
- 2. Create several implementations of RetryPolicy and provide configuration properties for them.
- 3. Implement heath checks by sending regular requests to the /status endpoint of configured servers. Add configuration properties so that you can customize endpoint value

#### Caching





#### **Caffeine**



- ✓ High-performance caching library based on Java 8
- ✓ Inspired by Guava-based cache and ConurrentHashMap
- ✓ Asynchronous cache loading
- ✓ Sized-base eviction
- ✓ Time-based expiration
- Eviction notification
- Writes propagation
- Cache access statistics
- ✓ JCache support

#### Caffeine. Examples



```
Cache<Integer, Book> cache = Caffeine.newBuilder()
        .expireAfterWrite(10, TimeUnit.MINUTES)
        .maximumSize(10_000)
        .build();
    Book book = cache.getIfPresent(1);
    Book book2 = cache.get(2, id -> new Book(id));
    cache.put(3, new Book(3));
    cache.invalidate(3);
                                           Calculate value if
                                            key is absent
    Removes key from cache Insert or update key
```

#### **Caffeine API**



Annotation	Description
initialCapacity()	Setups initial capacity of the cache
maximumSize()	Maximum amount of entries in the cache
expiresAfterWrite()	Specifies duration after the last insert(update) when entry should removed
expiresAfterAccess()	Specifies duration after the last read(write) when entry should removed from cache
removalListener()	Specifies listener after entry removal
recordStats()	Enables accumulation of statistics

#### **Cache API**



Annotation	Description
getIfPresent()	Returns key value of null
get()	Returns of compute cache value
put()	Inserts or updates key value
invalidate()	Removes cache entry
invalidateAll()	Discards all the entries in the cache
stats()	Returns current snapshot of the cache statistics
cleanUp()	Performs pending maintenance operations

#### Caffeine. Maven dependencies



```
<dependency>
     <groupId>com.github.ben-manes.caffeine</groupId>
          <artifactId>caffeine</artifactId>
                <version>2.6.2</version>
</dependency>
```

#### Task 9. Caching



- 1. Add Caffeine dependency to the order-service sub-project:
- Update load-balancer component to cache responses of the book-service. Assume that books objects are immutable and never change.
- 3. Add configuration properties that store cache properties: maximum and initial size, expiration.
- 4. Run integration tests and verify that orderController works properly.



#### Server load balancer

#### Task 10. Starting server balancer



- 1. Review server-balancer sub-project.
- 2. Import client-balancer library into server-balancer subproject.
- 3. Add necessary endpoints for server-balancer project
- 4. Update order-service project so that it uses server-side load balancing.



#### **Caching**



- ✓ Internal implementation by Spring
- ✓ JCache (JSR-107) is supported
- ✓ Various 3<sup>rd</sup> party implementations (EHCache, HazelCast, Infinispan, Couchbase, Redis, Caffeine)



#### Caching. Maven dependencies



```
<dependency>
     <groupId>com.github.ben-manes.caffeine</groupId>
          <artifactId>caffeine</artifactId>
                <version>2.6.2</version>
</dependency>
```

## **Spring annotations**



Name	Description
@EnableCaching	Enables annotation-driven cache management
@CacheConfig	Allows to config cache parameters
@Cacheable	Indicate that method(or all methods) result should be cached depending on the method arguments
@CachePut	Indicate that method result should be put into cache whereas method should be always invoked
@CacheEvict	Indicates that specific (or all) entries in the cache should be removed

#### **Enable caching**



```
@SpringBootApplication
@EnableCaching
@CacheConfig(cacheNames= {"orders", "payments"})
public class RestApplication {
                                         Optionally specify
                                         cache names
@Cacheable("books")
@GetMapping(path = "/{id}")
public Book findById(@PathVariable int id) {
     return bookRepository.findById(id);
 }
                                     Cache returned value
```

#### **Update cache**

```
DISCOVERY
```

Calls method and cache returned value

```
@DeleteMapping("/{id}")
@CacheEvict("books")
public void deleteBook(@PathVariable int id) {
    bookRepository.delete(id);
}
Calls method and
remove cache entry
```

#### **Advanced caching**



```
@GetMapping("/{id}")
@Cacheable(value = "books", key = "#id")
public Book findById(@PathVariable String\id) {
    return bookService.findBook(id);
                                            Specify cache
                                             kev
@PutMapping("/{id}")
@CachePut(value = "books", key = "#id")
public Book update(@PathVariable int id,
        @RequestBody Book book) {
```

```
spring.cache.type=caffeine ← Specify exact cache provider
```

### Task 11. Spring Cache



- Add @EnableCaching annotation to the LoadBalancerApplication class.
- 2. Add Cache dependency
- 3. Add @Cacheable annotation to the GET REST-services. How does it affect its behavior?
- 4. Add @CachePut annotation on POST/PUT REST services and @CacheEvict annotation on DELETE REST services. Verify its behavior.



## **Authentication**

Из диалога двух программистов:

- Кажется, у нас дыра в безопасности!
- Слава Богу, хоть что-то у нас в безопасности...

## **Authentication**

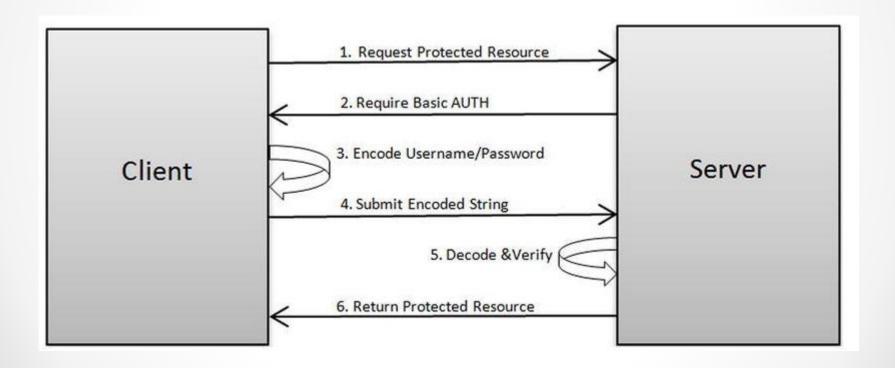


- ✓ Basic
- Digest
- ✓ Token (JWT)
- ✓ Digital Signature
- ✓ Certificate
- ✓ OAuth 1.0/2.0



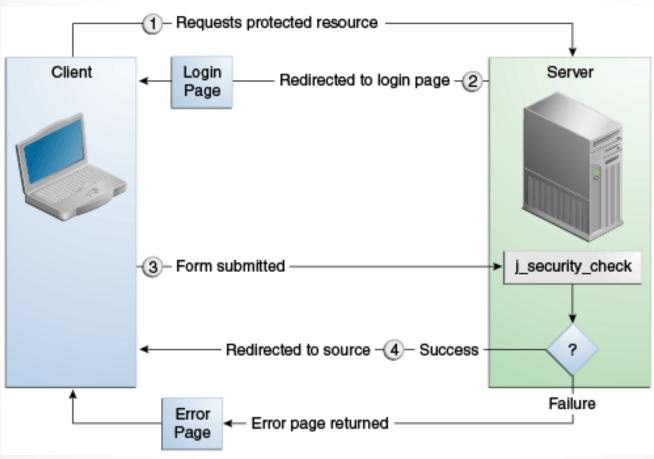
## **Basic authentication**





## Form-based authentication





# **Spring Security**



- ✓ Formerly Acegi Security
- ✓ Support for authentication and authorization
- ✓ Integration with Spring MVC(REST)
- ✓ Attacks protections (CSRF, session fixation)
- ✓ Basic, Digest, CAS, OpenID, JAAS and LDAP authentication
- ✓ Integration testing
- ✓ Based on filter chain



# **Spring Security. Maven**



```
<dependency>
     <groupId>org.springframework.boot</groupId>
     <artifactId>spring-boot-starter-security</artifactId>
     <version>${spring.boot.version}</version>
</dependency>
```

# **Configuration steps**



- ✓ Declare protected resources
- ✓ Setup authentication
- ✓ Setup authorization (optional)
- ✓ Manage user storage

# **Spring Security**



```
@Configuration
@EnableGlobalMethodSecurity(prePostEnabled = true)
public class SecurityConfiguration extends
                WebSecurityConfigurerAdapter {
    @Override
    protected void configure(HttpSecurity http)
            throws Exception {
        http.authorizeRequests().anyRequest()
            .fullyAuthenticated();
        http.httpBasic();
        http.csrf().disable();
```

Enable POST/PUT requests

# **Spring Security. Users**



```
@Bean
public PasswordEncoder passwordEncoder() {
    return NoOpPasswordEncoder.getInstance();
}

Can be BCrypt, SCrypt, MD4
```

```
spring.security.user.name=admin
spring.security.user.password=admin
spring.security.user.roles=USER,ADMIN
```

application.properties

## **Spring Security. Users**



```
@Override
protected void configure(AuthenticationManagerBuilder auth)
    throws Exception {
        auth.inMemoryAuthentication()
            .withUser("john").password("123").roles("USER")
            .and()
            .withUser("donald").password("abc").roles("ADMIN");
}
```

class SecurityConfiguration

# **Rest Client. Supply credentials**

```
public class RestClient {
    private TestRestTemplate restTemplate;

    public RestClient() {
        restTemplate = new TestRestTemplate("admin", "admin");
    }
}
```

DISCOVERY

## **Unit-testing**



## **Unit-testing**



```
@Test
@WithMockUser(username = "user", authorities = { "USER" })
public void findBooks_StorageIsNotEmpty_OneBookReturned()
        throws Exception {
    given(bookRepository.findAll()).
            willReturn(Arrays.asList(new Book()));
    mockMvc.perform(get("/book")).andExpect(status().isOk())
            .andExpect(content().contentType(
                    MediaType. APPLICATION JSON_UTF8 VALUE))
            .andExpect(jsonPath("$", Matchers.hasSize(1)));
```

#### **Task 12. Spring Security**



- 1. Add Spring Security dependency to server-balancer project
- 2. Add Spring Security Test dependencies
- 3. Add **Spring Security** configuration:
  - 3.1 Add user(s) with different roles
  - 3.2 Add HTTP basic authentication
- 4. Update order-service project and load-balancer component so that it sends basic authentication credentials.

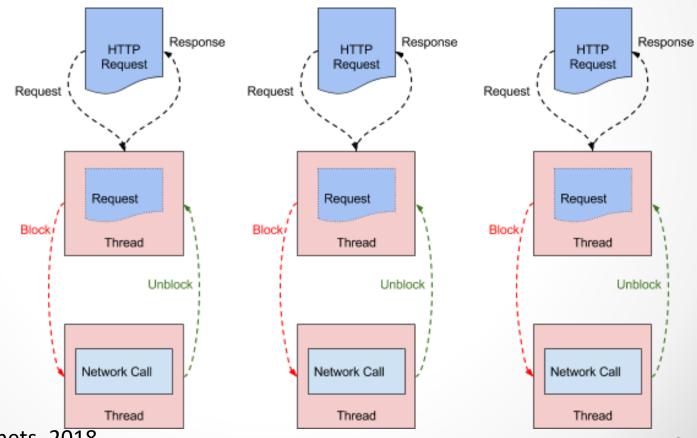






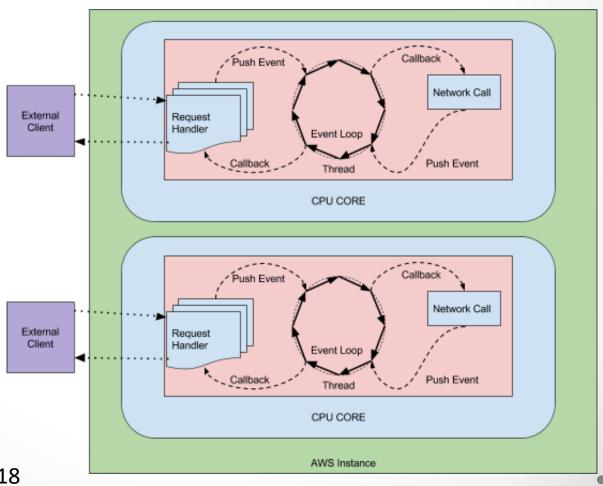
## Blocking I/O





## Non-blocking I/O





## **Spring 5 WebFlux**



- New spring-web-flux module adapting Reactive Streams specification
- Based on Project Reactor
- Reusing Spring MVC programming model but in non-blocking mode
- Based on non-blocking Servlet API or native SPI connectors(Netty, Undertow)
- Supported by Tomcat/Jetty/Netty/Undertow
- Reactive client WebClient
- Support unit-testing with WebTestClient

### **Spring Web Flux**



@Controller, @RequestMapping

**Router Functions** 

spring-webmvc

spring-webflux

Servlet API

HTTP / Reactive Streams

Servlet Container

Tomcat, Jetty, Netty, Undertow

### **Spring Web Flux**



```
@FunctionalInterface
public interface HandlerFunction<T extends ServerResponse> {
    /**
    * Handle the given request.
    * @param request the request to handle
    * @return the response
    */
    Mono<T> handle(ServerRequest request);
```



#### SPRING INITIALIZR bootstrap your application now

Generate a Maven Project with Java	and Spring Boot 2.0.0 RC2
Project Metadata	Dependencies
Artifact coordinates	Add Spring Boot Starters and dependencies to your application
Group	Search for dependencies
com.example	reactive web
Artifact	Reactive Web
demo	Reactive web development with Netty and Spring WebFlux
Generate Project alt + ₪	

### Spring Web Flux. Maven



```
<dependency>
     <groupId>org.springframework.boot</groupId>
     <artifactId>spring-boot-starter-webflux</artifactId>
     <version>${spring.boot.version}</version>
</dependency>
```

#### **REST** controller. Mono



```
@RestController
@RequestMapping
public class ProductController {
    private ProductService productService;
    @GetMapping(path = "/{id}")
    public Mono<Product> get(
            @PathVariable("id") int id) {
        return productService.find(id);
```

#### **REST** controller. Flux



```
@RestController
@RequestMapping("/random")
public class RandomController {
    private Random random = new Random();
    @GetMapping
    public Flux<Integer> generate() {
        return Flux.fromStream(Stream.
                generate(random::nextInt)).
                delayElements(Duration.ofSeconds(1));
```

#### **REST** controller. Flux



```
@RestController
@RequestMapping("/random")
public class RandomController {
    private Random random = new Random();
    @GetMapping(produces=MediaType.TEXT_EVENT_STREAM_VALUE)
    public Flux<Integer> generate() {
        return Flux.fromStream(Stream.
                generate(random::nextInt)).
                delayElements(Duration.ofSeconds(1));
```

#### Web client



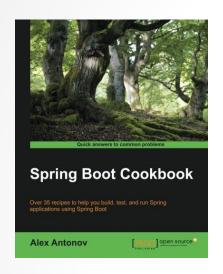
### Task 13. Spring Web Flux

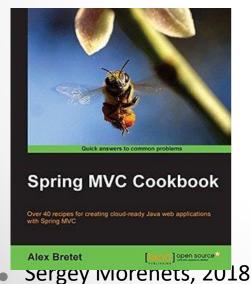


- 1. Add Sping WebFlux dependency to pom.xml in orderservice/server-balancer projects:
- 2. Update REST services in book-service or server-balancer applications so that uses **Spring Web Flux** API
- 3. Update order-service application and balancer component and switch from RestTemplate to **WebClient** as REST client.

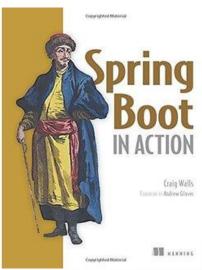


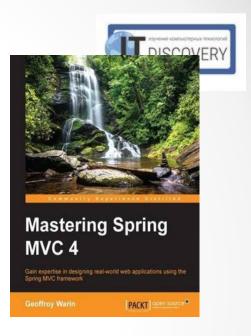
### **Books**

















- ✓ Sergey Morenets, <a href="mailto:sergey.morenets@gmail.com">sergey.morenets@gmail.com</a>
- Sergey Morenets, 2018